

forest_fires

November 9, 2020

```
[144]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('ggplot')

from scipy.stats import zscore

[145]: # read in forest fires dataset into a pandas dataframe
df = pd.read_csv('forest-fires.csv')

# perform very basic analysis of the metadata (number of points and data types
→for each column)
print(df.shape)
print(df.dtypes)
df.describe().T
```

(517, 13)

```
x          int64
y          int64
month      object
day        object
ffmc       float64
dmc        float64
dc         float64
isi        float64
temp       float64
rh         int64
wind       float64
rain       float64
area       float64
dtype: object
```

```
[145]:
```

	count	mean	std	min	25%	50%	75%	max
x	517.0	4.669246	2.313778	1.0	3.0	4.00	7.00	9.00
y	517.0	4.299807	1.229900	2.0	4.0	4.00	5.00	9.00
ffmc	517.0	90.644681	5.520111	18.7	90.2	91.60	92.90	96.20
dmc	517.0	110.872340	64.046482	1.1	68.6	108.30	142.40	291.30

dc	517.0	547.940039	248.066192	7.9	437.7	664.20	713.90	860.60
isi	517.0	9.021663	4.559477	0.0	6.5	8.40	10.80	56.10
temp	517.0	18.889168	5.806625	2.2	15.5	19.30	22.80	33.30
rh	517.0	44.288201	16.317469	15.0	33.0	42.00	53.00	100.00
wind	517.0	4.017602	1.791653	0.4	2.7	4.00	4.90	9.40
rain	517.0	0.021663	0.295959	0.0	0.0	0.00	0.00	6.40
area	517.0	12.847292	63.655818	0.0	0.0	0.52	6.57	1090.84

```
[146]: # Inspect the first few data points to gain a brief understanding of the data
df.head()
```

```
[146]:   x  y month  day  ffmc  dmc   dc  isi  temp  rh  wind  rain  area
0  7  5  mar  fri  86.2  26.2  94.3  5.1   8.2  51   6.7   0.0   0.0
1  7  4  oct  tue  90.6  35.4  669.1  6.7  18.0  33   0.9   0.0   0.0
2  7  4  oct  sat  90.6  43.7  686.9  6.7  14.6  33   1.3   0.0   0.0
3  8  6  mar  fri  91.7  33.3   77.5  9.0   8.3  97   4.0   0.2   0.0
4  8  6  mar  sun  89.3  51.3  102.2  9.6  11.4  99   1.8   0.0   0.0
```

```
[147]: # Find missing values and correct them in the dataset if needed
print(df.isna().sum().sum())
```

0

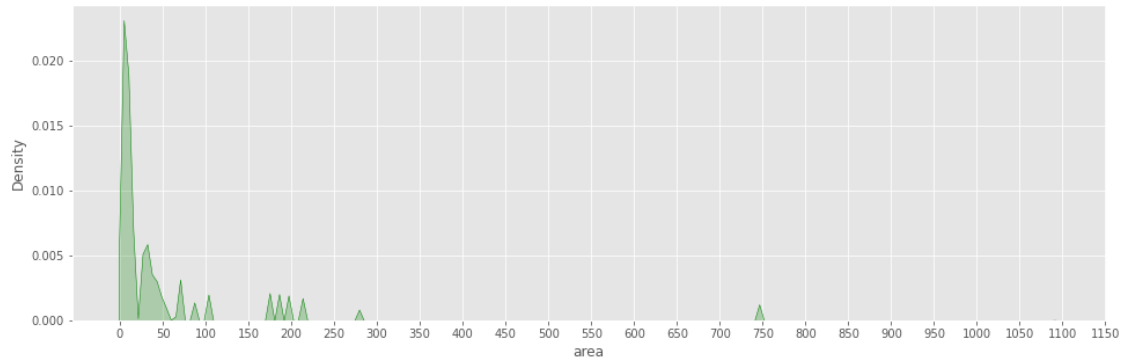
```
[148]: # Configure plotting library
plt.rcParams['figure.figsize'] = 9,5
```

```
[149]: # Analyze skew and kurtosis values
## Skewness: A measure of how skewed (non-symmetric) the data is relative to
→the midpoint.
print("Area skewness: {}".format(df['area'].skew()))
## Kurtosis: A measure of how heavy the tails are in the data (how many
→outliers there are)
print("Area kurtosis: {}".format(df['area'].kurtosis()))
```

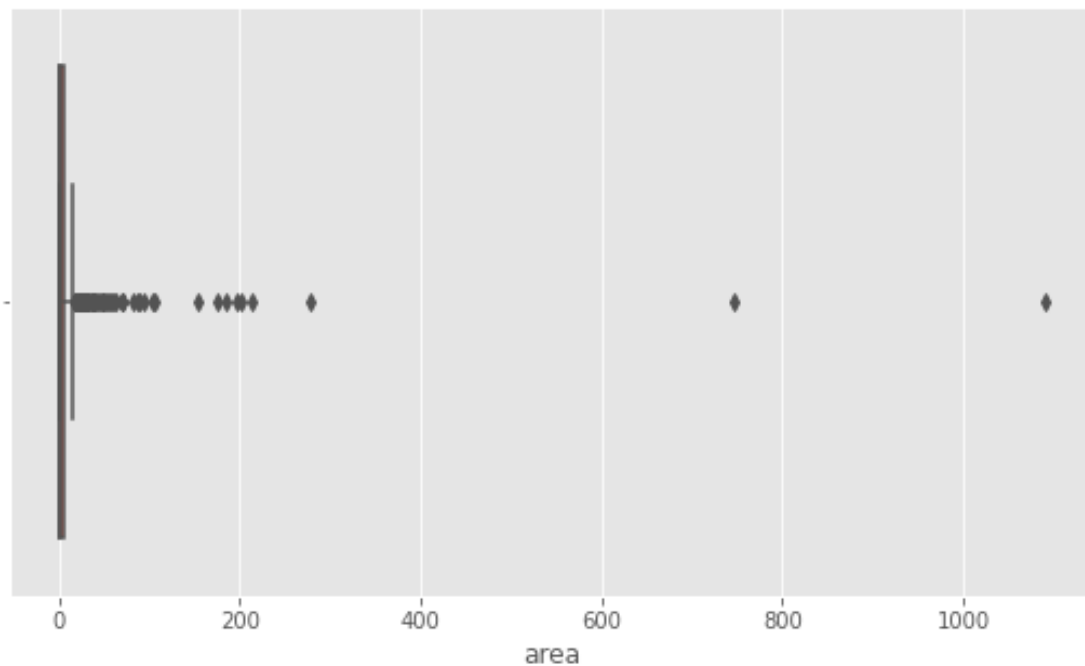
Area skewness: 12.846933533934868

Area kurtosis: 194.1407210942299

```
[150]: # Plot a Kernel Density Estimate of our data. This is essentially a histogram
→but it provides more useful insights.
plt.figure(figsize=(16,5))
ax = sns.kdeplot(df['area'],bw_adjust=0.02,shade=True,color='g')
plt.xticks([i for i in range(0,1200,50)])
plt.show()
```



```
[151]: # Plot a boxplot of all the area values to find any outliers
ax = sns.boxplot(x=df['area'])
```



```
[152]: # Notes:
## The data is very skewed by the outliers.
## We see that most forest fires cover less than 50 hectares of land.
## The main outliers have been identified below
outliers = df[abs(zscore(df['area'])) >= 3 ]
outliers
```

```
[152]:
```

	x	y	month	day	ffmc	dmc	dc	isi	temp	rh	wind	rain	area
237	1	2	sep	tue	91.0	129.5	692.6	7.0	18.8	40	2.2	0.0	212.88
238	6	5	sep	sat	92.5	121.1	674.4	8.6	25.1	27	4.0	0.0	1090.84

415	8	6	aug	thu	94.8	222.4	698.6	13.9	27.5	27	4.9	0.0	746.28
479	7	4	jul	mon	89.2	103.9	431.6	6.4	22.6	57	4.9	0.0	278.53

```
[153]: # Create a new dataframe without the area (for which we will later create a
        ↳ prediction model)
df_input = df.drop(columns='area')

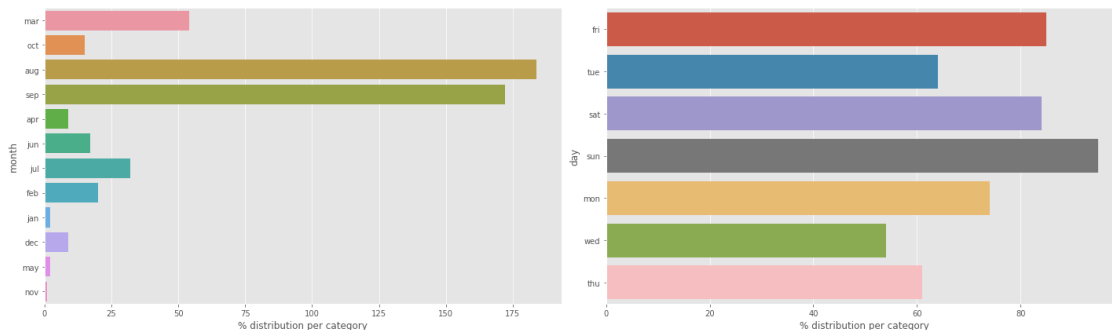
# Split up the data columns into categorical columns and numerical columns
categorical = df_input.select_dtypes(include='object').columns.tolist()
numerical = df_input.select_dtypes(exclude='object').columns.tolist()

print(categorical)
print(numerical)
```

```
['month', 'day']
```

```
['x', 'y', 'ffmc', 'dmc', 'dc', 'isi', 'temp', 'rh', 'wind', 'rain']
```

```
[154]: # Visualize categorical data
plt.figure(figsize=(20,6))
for index, cat_column in enumerate(categorical, start=1):
    plt.subplot(1,2,index)
    sns.countplot(data=df_input,y=cat_column)
    plt.ylabel(cat_column)
    plt.xlabel('% distribution per category')
plt.tight_layout()
plt.show()
```



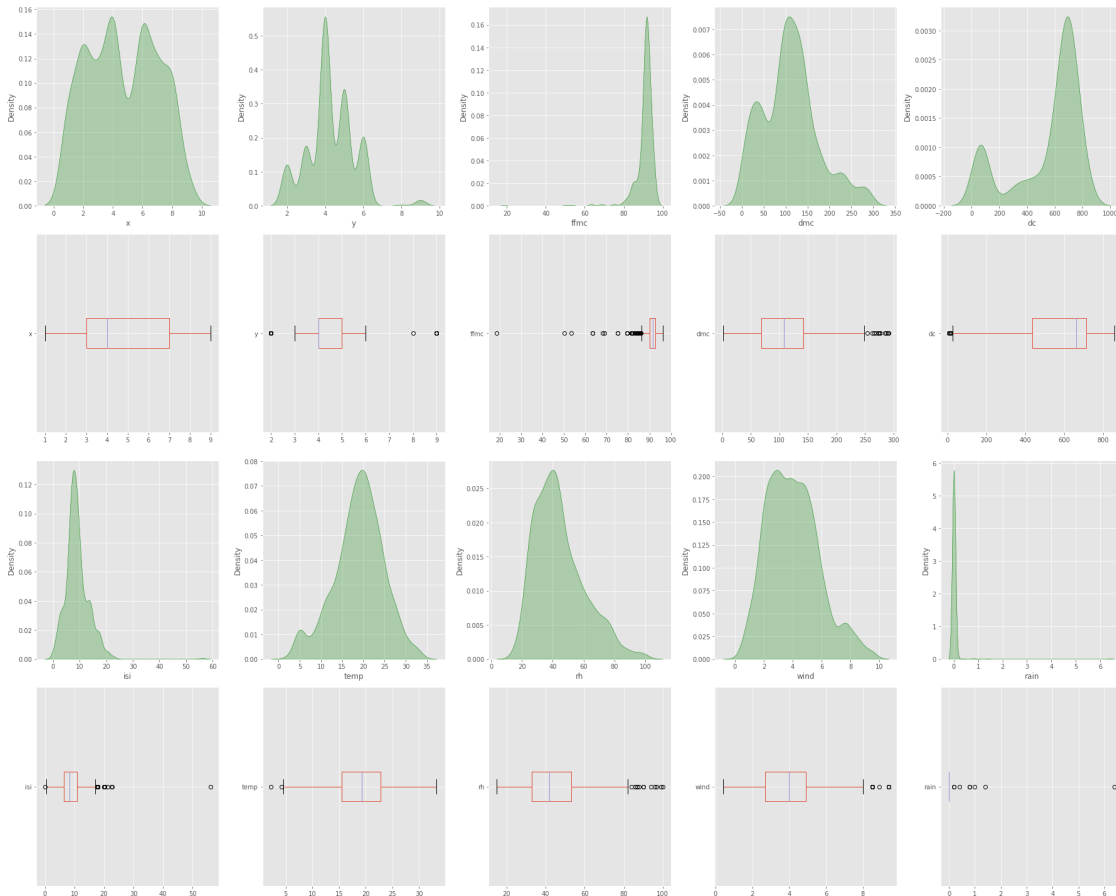
```
[155]: # We notice a huge spike in forest fires during August and September.
        # Sunday seems to have more fires on average, but there seems to be little
        ↳ significant difference.
```

```
[156]: # Visualize numerical data
plt.figure(figsize=(25,20))
for index, num_column in enumerate(numerical, start=1):
    plot_index = index+5 if index > 5 else index
    plt.subplot(4,5,plot_index)
```

```

sns.kdeplot(data=df[num_column],bw_adjust=0.8,color='g',shade=True)
plt.subplot(4,5,plot_index+5)
df[num_column].plot.box(vert=False)
plt.tight_layout()
plt.show()
numerical_data = df[numerical]
pd.DataFrame(data=[numerical_data.skew(),numerical_data.
→kurtosis()],index=['skewness','kurtosis'])

```



[156]:

	x	y	ffmc	...	rh	wind	rain
skewness	0.036246	0.417296	-6.575606	...	0.862904	0.571001	19.816344
kurtosis	-1.172331	1.420553	67.066041	...	0.438183	0.054324	421.295964

[2 rows x 10 columns]

[157]: # Notes:
The KDE plots and box plots reveal that there is the most skew and kurtosis
→ in FFMC, ISI, and rain
These columns also had the most extreme outliers.

```
[158]: # Create categorical representations of the damage from forest fires based on
→ area
def damage_level(area):
    if area == 0.0:
        return "No damage"
    elif area <= 1:
        return "low"
    elif area <= 25:
        return "moderate"
    elif area <= 100:
        return "high"
    else:
        return "very high"

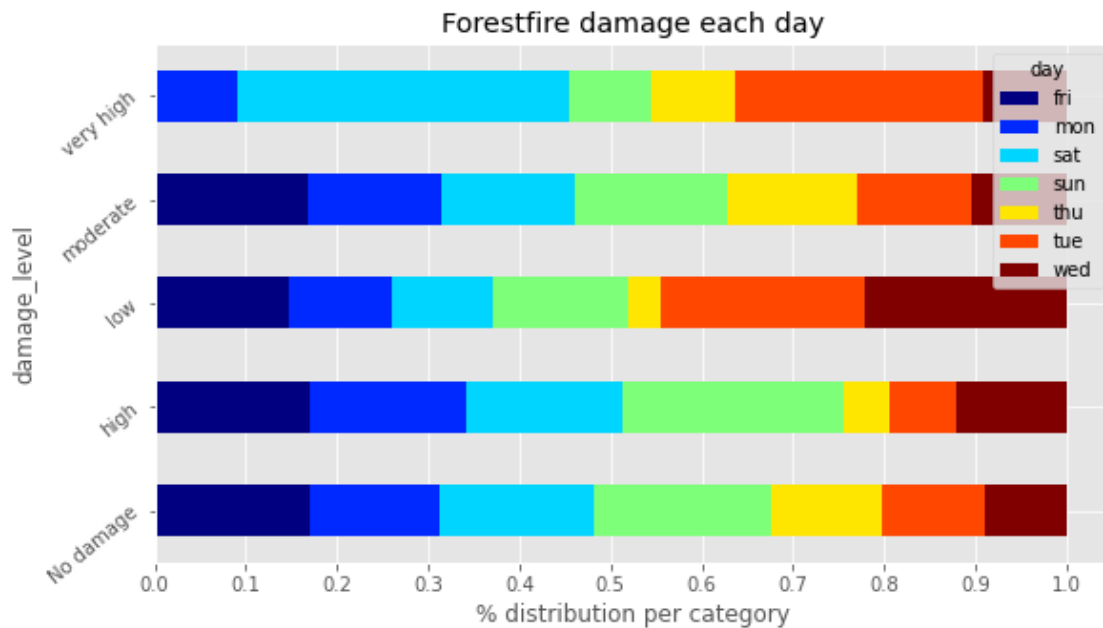
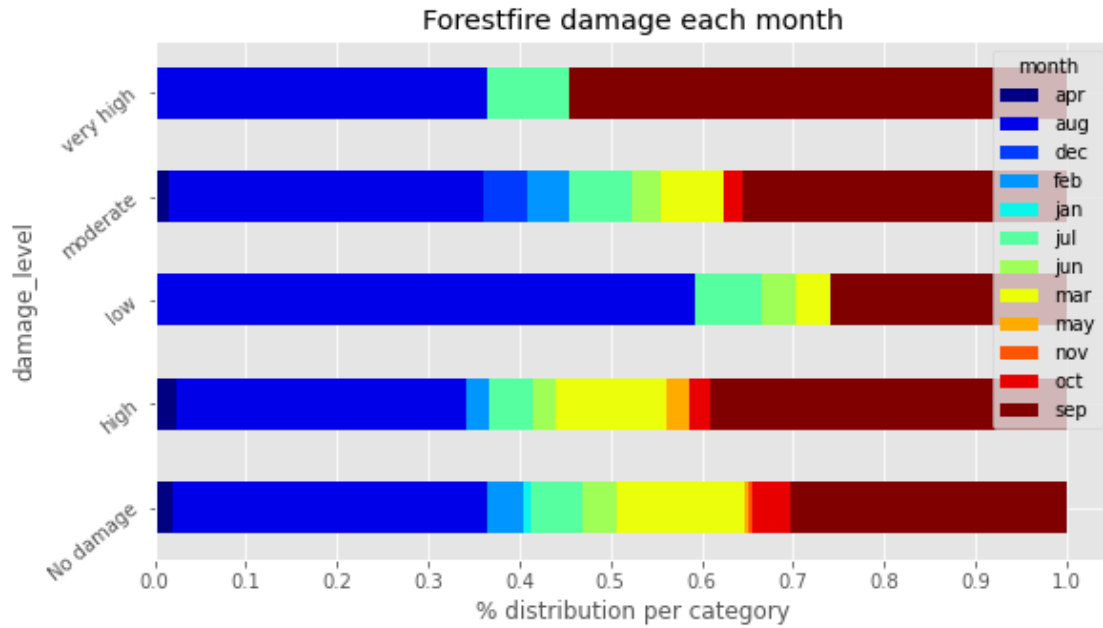
df['damage_level'] = df['area'].apply(damage_level)
df.head()
```

```
[158]:
```

	x	y	month	day	ffmc	dmc	...	temp	rh	wind	rain	area	damage_level
0	7	5	mar	fri	86.2	26.2	...	8.2	51	6.7	0.0	0.0	No damage
1	7	4	oct	tue	90.6	35.4	...	18.0	33	0.9	0.0	0.0	No damage
2	7	4	oct	sat	90.6	43.7	...	14.6	33	1.3	0.0	0.0	No damage
3	8	6	mar	fri	91.7	33.3	...	8.3	97	4.0	0.2	0.0	No damage
4	8	6	mar	sun	89.3	51.3	...	11.4	99	1.8	0.0	0.0	No damage

[5 rows x 14 columns]

```
[159]: # Visualize the damage level with respect to the month and day of the event
for index, category in enumerate(categorical, start=1):
    cross = pd.
→ crosstab(index=df['damage_level'], columns=df[category], normalize='index')
    cross.plot.barh(stacked=True, rot=40, cmap='jet')
    plt.xlabel('% distribution per category')
    plt.xticks(np.arange(0,1.1,0.1))
    plt.title("Forestfire damage each {}".format(category))
plt.show()
```



[160]: # Notes:
 ## This reveals that although august and september had the most fires, many of them were relatively low damage.
 ## Furthermore, we can notice that every high damage fire happened within three months: July, August, September.

```
## May, April, November, and December had almost no fires, which coincides with  
→our previous visualization that didn't factor in damage done.  
## The by day graph provides little valueable insight, since almost every day  
→has an even amount of each damage level.
```

```
[161]: # Visualize the damage level with respect to the numerical columns of the data  
plt.figure(figsize=(20,40))  
for index, category in enumerate(numerical, start=1):  
    plt.subplot(10,1,index)  
    if category in ['X','Y']:  
        sns.swarmplot(data=df,x=category,y='area',hue='damage_level')  
    else:  
        sns.scatterplot(data=df,x=category,y='area',hue='damage_level')  
plt.show()
```